

Inter-Process Communication

CPEN333 – Software Design for Engineers II
2023 W1
University of British Columbia

© *Farshid Agharebparast*



Introduction

- We have already discussed cooperating processes.
- Cooperating processes need **interprocess communication (IPC)** that would allow them to exchange data and information.
- In this set of slides, we examine different mode of communications and methods.

Objectives

- To describe the communication within processes
 - ❖ message passing
 - ❖ shared-memory
- To describe communication in client-server systems

Cooperating processes


- Reasons for cooperating processes could be:
 - ❖ **Information sharing**: e.g. several users may be interested in the same piece of information
 - ❖ **Computation speedup**: e.g. breaking a particular task into subtasks being executed in parallel
 - ❖ **Modularity**: e.g. dividing a system functions into separate processes in a modular fashion
 - ❖ **Convenience**: even an individual user may work on many tasks at the same time

Communications Models

- Two main models of IPC
 - ❖ **Shared memory**: a region of memory is used that is shared by the cooperating processes
 - ❖ **Message passing**: Communication takes place by means of messages exchanged between cooperating processes
- Both of these two models are common in the OSs, and many systems implement both.

Shared Memory Systems

- Recall that, normally, the OS tries to prevent one process from accessing another process's memory.
 - ❖ **Shared-memory** requires that two or more processes agree to remove this restriction.

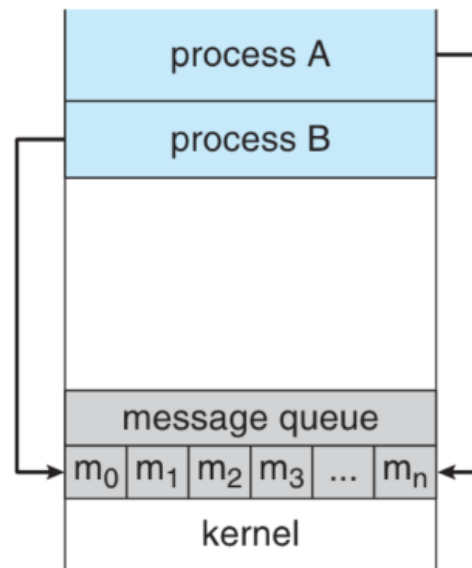


Shared Memory (cont.)

- Usually the shared-memory resides in the address space of the process creating the shared-memory segment
- Then those processes can exchange information by reading and writing data in the shared areas (not under the OS's control)
- **Synchronization** may be needed, of course: e.g., the processes are responsible for ensuring that they are **not** writing to the same location simultaneously.
 - ❖ We have already seen this as an example in the **producer-consumer problem**.


Message Passing Systems

- **Message passing** is the other method that provides a mechanism for processes to communicate and to synchronize their actions
 - ❖ Processes communicate with each other without resorting to shared variables
 - ❖ The operating system is to provide the means for cooperating processes to communicate with each other via a message-passing facility.




Message Passing (cont.)

- A general message passing facility provides at least two operations:
 - ❖ `send(message)` and `receive(message)`
- If two processes P and Q wish to communicate, they need to:
 - ❖ establish a *communication link* between them
 - ❖ exchange messages via send/receive
- A particularly useful and practical method in distributed environment (e.g. chat programs)



P and Q on two different machines



P and Q on the same machine

Implementation Questions


- So far, we have discussed general message passing concepts.
- For the implementation, we might decide on several options.
- How are links established?
 - ❖ A **communication link** must exist between the cooperating processes.
 - ❖ The link can be
 - **physical**, e.g. a hardware bus, a communication network (LAN or Internet)
 - **logical**, e.g. a message passing queue
- Can a link be associated with more than two processes?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Blocking vs non-blocking

- Communication between processes takes place through calls to `send()` and `receive()` primitives
- There are different design options for implementing each primitive:
 - ❖ Message passing may be either **blocking** or **non-blocking**
- Blocking is considered *synchronous*
 - ❖ **Blocking send** has the sender block until the message is received
 - ❖ **Blocking receive** has the receiver block until a message is available
- Non-blocking is considered *asynchronous*
 - ❖ **Non-blocking send** has the sender send the message and continue
 - ❖ **Non-blocking receive** has the receiver receive a valid message or a null
- Different combination of `send()` and `receive()` are possible. When Both send and receive are blocking, we have a rendezvous between them.


Buffering

- Whether communication is direct or indirect, messages exchanged by communicating processes reside in a temporary queue.
- Such queues can be implemented in one of three ways
 - ❖ **Zero capacity** (queue max length is zero, i.e., 0 messages)
 - Sender must wait for the receiver
 - ❖ **Bounded capacity** (finite length of n messages)
 - Sender must block if the link is full, otherwise it can continue without waiting
 - ❖ **Unbounded capacity** (infinite length)
 - Sender never waits



Communications in Client-Server Systems

- So far, it was described how processes can communicate using:
 - ❖ shared memory and
 - ❖ message passing
- Strategies for communication in client-server systems:
 - ❖ Sockets
 - ❖ Pipes
 - ❖ Remote Procedure Calls (RPC)

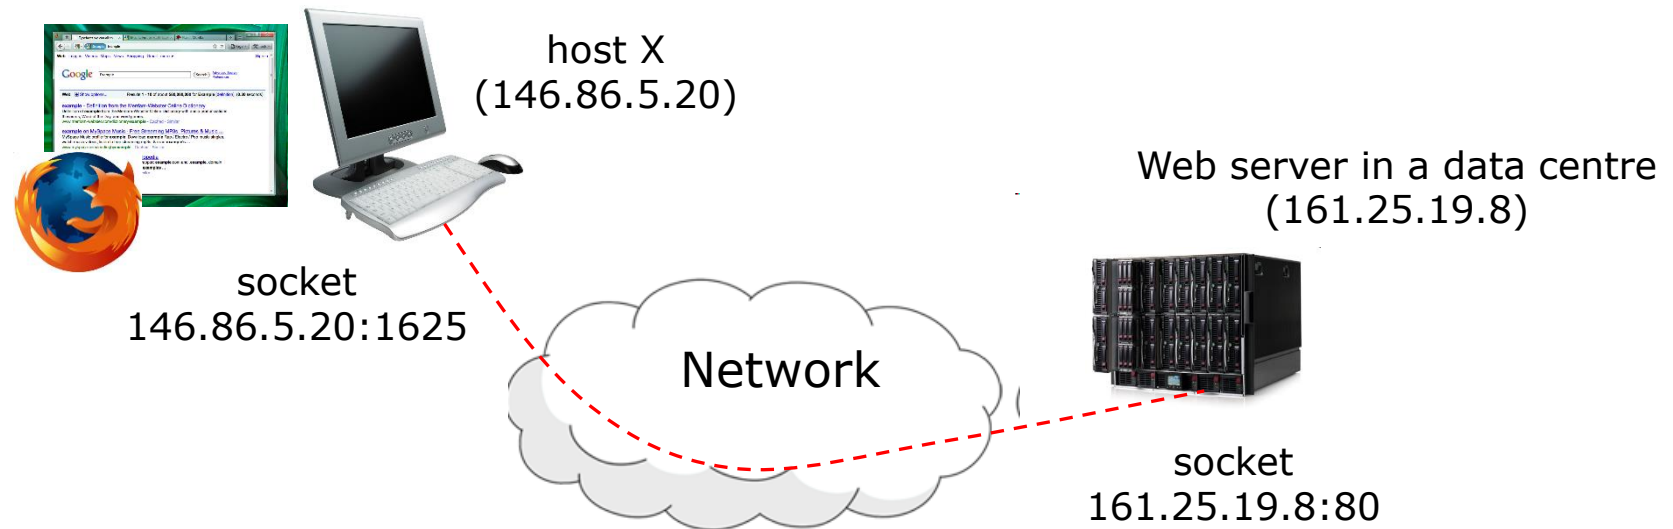


Client-server model

- The primary model used in the Internet: email, web, ...
- The client-server describes the relationship of the cooperating processes.
 - ❖ A **server** is the service provider, and generally is constantly awaiting to receive incoming requests to provide service.
 - A server is described by the service it provides, e.g. a file server or a web server.
 - ❖ **Clients** are service requesters and initiate communication sessions with the server to receive service.
 - Similarly we have email client, web client (e.g. a web browser), ...

Sockets

- A **socket** is defined as an *endpoint for communication*
 - ❖ It is identified by an IP address concatenated with a port number, in its simplest form.
 - e.g., the socket **146.86.5.20:1625** refers to port **1625** on host **146.86.5.20**
- A pair of processes communicating over network employ a pair of sockets.



TCP vs UDP sockets

- Two very popular types of sockets are TCP and UDP.
 - ❖ **TCP** (transmission control protocol) and **UDP** (user datagram protocol) are the Internet's primary transport protocols.
- TCP sockets are *connection-oriented*, that is a connection is first established and then data can transfer.
- UDP sockets are *connectionless*, that is, no connection is established first. Each data segment is individually addressed and routed to be sent.
- UDP and TCP sockets are different in some other respects, discussed later.


Pipes

- A **pipe** acts as a conduit providing one of the simpler ways for processes to communicate
 - ❖ Ordinary pipes allow two processes to communicate in standard producer-consumer fashion
 - The producer writes to one end of the pipe (write-end) and the consumer reads from the other end (read-end)



- A pipe creates an (r, w) file descriptor: r is for reading and w for writing.

- ❖ For example, if the parent process wants to receive data from the child process, it keeps r open, and the child keeps w open



Pipes (cont.)

- Pipes are used quite often in the command-line environment (originally from the UNIX OS) in which the output of one command serves as input to the second.
- A pipe can be constructed on the CLI using the | character
 - ❖ Example (macOS or Linux terminal): `ls | more`
 - ❖ Equivalent example (Windows cmd/powershell): `dir | more`

Remote Procedure Calls

- **Remote procedure call** (RPC) abstracts procedure calls between processes on networked systems
- The semantics of RPCs allow a client to invoke a procedure on a remote host as it would invoke a procedure locally
- We must use a message-based communication to provide remote service
 - ❖ The messages are well structured
 - ❖ Each message is addressed to an RPC daemon (e.g. background server) listening to a port on a remote system, and contains
 - the identifier of the function to execute and
 - the parameters to pass to that function

Remote Procedure Calls (cont.)

- An important issue that must be dealt with concerns differences in data representation of the client and server machines.
- **Endianness** describes the order of bytes of a word in a computer memory.
 - ❖ **Little-endian**: some systems store the least significant byte first.
 - e.g. Intel's x86 processors and their clones are little endian.
 - ❖ **Big-endian**: some systems store the most significant byte first.
 - e.g. Motorola 6800, PowerPC
 - ❖ **Bi-endian**: allows switchable endianness
 - e.g. ARM architecture since version 3 (little-endian generally)

References

- Some sections of chapter 3 of Operating Systems Concepts

Acknowledgement: This set of slides is partly based on the PPTs provided by the Wiley's companion website for the operating system concepts book (including textbook images, when not explicitly mentioned/referenced).

